

```

In [1]: import numpy as np
        from collections import defaultdict

        # Step 1: Read files
        '''
        Rewards = 81-d column array
        Probs = dict{action : {s:[(s1', p1),...]}]}
        '''

        REWARD_FILE = 'rewards.txt'
        P_FILE = 'prob_a.txt'

        def read_rewards(fname):
            with open(REWARD_FILE) as f:
                return np.mat([int(line.strip()) for line in f]).T

        def read_probs(fname):
            probs = dict()
            for action in range(1, 5):
                probs[action] = defaultdict(list)
                name, ext = fname.split('.')
                filename = name + str(action) + '.' + ext
                with open(filename) as f:
                    for line in f:
                        content = line.split()
                        s_cur, s_next, p = int(content[0]), int(content[1]), float(content[2])
            ])
                probs[action][s_cur].append((s_next, p))
            return probs

        rewards = read_rewards(REWARD_FILE)
        probs = read_probs(P_FILE)

        # Step 2: Initialize variables
        gamma = 0.99
        states = range(1, 82)
        actions = range(1, 5)
        policy = {state: np.random.randint(low = 1, high = 5) for state in states}

        # Step 3: Start policy iteration

        def evaluate_values():
            M = np.eye(len(states))
            for state in states:
                action = policy[state]
                for s_next, p in probs[action][state]:
                    M[state - 1, s_next - 1] -= gamma * p
            return np.linalg.solve(M, rewards)

        # values = np.mat() : MINUS 1
        # Evaluate Values
        def q_sa(state, action, values):
            reward = 0
            for s_next, p in probs[action][state]:
                reward += p * values[s_next - 1]
            return rewards[state - 1] + gamma * reward

        # Greedy update policy
        def update_policy(values):
            is_updated = False
            policy_new = {state: None for state in states}
            for state in states:
                q_max, action_best = float('-inf'), None

```

```

    for action in actions:
        q_sa_value = q_sa(state, action, values)
        if q_max < q_sa_value:
            q_max, action_best = q_sa_value, action
    policy_new[state] = action_best
    if action_best != policy[state]:
        is_updated = True
    return is_updated, policy_new

is_updated = True
iter = 0
while is_updated:
    values = evaluate_values()
    is_updated, policy = update_policy(values)
    iter += 1
print("Iteration = {}".format(iter))

best_value = evaluate_values().reshape((9, 9)).T
print(best_value)
# np.savetxt('bestvalue.txt', best_value, fmt='%g')
best_policy = np.array([action for _, action in sorted(list(policy.items()))]).reshape((9, 9)).T
print(best_policy)
# np.savetxt('bestpolicy.txt', best_policy, fmt='%g')

```

Iteration = 4

```

[[ 0. 0. 0. 0. 0.
  0. 0. 0. 0. ]
 [ 0. 65.77308407 67.13647421 77.84605 79.84451583
 72.47511769 -100. 0. 100. ]
 [ 0. 55.88294346 -100. 70.30818136 81.34440225
 83.04847989 84.88054612 96.87232244 98.71875987]
 [ 0. 54.92298013 50.47656297 59.66641187 0.
 80.95826449 0. 97.04482865 98.72729893]
 [ 53.50968756 54.14557214 0. -100. -100.
 61.77980767 -100. 88.22035599 100. ]
 [ 0. 52.50402036 43.9359876 51.09137525 61.00715483
 71.78642614 73.94661407 85.18458536 97.57257319]
 [ 0. 43.77254574 -100. 0. 0.
 70.35142939 0. -100. 88.40593622]
 [ 0. 47.95296148 48.76871928 58.14735126 59.39003194
 60.1688947 -100. 0. 100. ]
 [ 0. 0. 0. 0. 0.
 0. 0. 0. 0. ]]

[[1 1 1 1 1 1 1 1 1]
 [1 3 3 3 4 4 1 1 1]
 [1 2 1 3 3 3 3 3 2]
 [1 2 3 2 1 2 1 3 4]
 [3 2 1 1 1 2 1 3 1]
 [1 2 1 3 3 3 3 3 2]
 [1 2 1 1 1 2 1 1 4]
 [1 3 3 3 3 2 1 1 1]
 [1 1 1 1 1 1 1 1 1]]

```

```

In [2]: %matplotlib inline
import math
import numpy as np

import cv2
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
maze = mpimg.imread('Maze.jpg')
plt.rcParams['figure.figsize'] = [20, 20]

def draw_arrow(start, end, color, thickness, image, length = 20, alpha = 127):
    PI = 3.1415926
    angle = math.atan2(end[1] - start[1], end[0] - start[0])
    cv2.line(image, start, end, color, thickness)

    arrow_x = end[0] + length * math.cos(angle + PI * alpha / 180)
    arrow_y = end[1] + length * math.sin(angle + PI * alpha / 180)
    cv2.line(image, (int(arrow_x), int(arrow_y)), end, color, thickness)

    arrow_x = end[0] + length * math.cos(angle - PI * alpha / 180)
    arrow_y = end[1] + length * math.sin(angle - PI * alpha / 180)
    cv2.line(image, (int(arrow_x), int(arrow_y)), end, color, thickness)


def draw_direction(x, y, direction, image, color, thickness):
    width = 22
    side = 110
    x_center, y_center = x * side + width, y * side + width
    if direction == 1:
        start, end = (x_center + width, y_center), (x_center - width, y_center)
    elif direction == 2:
        start, end = (x_center, y_center + width), (x_center, y_center - width)
    elif direction == 3:
        start, end = (x_center - width, y_center), (x_center + width, y_center)
    elif direction == 4:
        start, end = (x_center, y_center - width), (x_center, y_center + width)

    biasx = biasy = 35
    start = (start[0] + biasx, start[1] + biasx)
    end = (end[0] + biasx, end[1] + biasx)
    draw_arrow(start, end, color, thickness, image)

numbered_square = set([5, 20, 22, 24, 26, 29, 30, 31, 33, 35, 38, 39, 42, 44, 57, 60,
75, 76, 78, 79] + list(range(66, 70)) + list(range(47, 54)) + list(range(11, 18)))
def is_numbered_square(x, y):
    place = x * 9 + y + 1
    return place in numbered_square

color = (255,0,0)
thickness = 10
arrow_image = np.copy(maze)
for x in range(9):
    for y in range(9):
        if is_numbered_square(x, y):
            draw_direction(x, y, best_policy[y, x], arrow_image, color, thickness)
plt.imshow(arrow_image)
plt.show()

```



```
In [3]: def value_iteration():
    v_t = np.array([0.0] * len(states))
    delta = float('inf')
    threshold = 0.0001
    while delta > threshold:
        delta = 0
        for state in states:
            best_value = float('-inf')
            for action in actions:
                best_value = max(best_value, rewards[state - 1] + gamma * sum(v_t[s_n
ext - 1] * p for s_next, p in probs[action][state]))
            delta = max(delta, abs(best_value - v_t[state - 1]))
            v_t[state - 1] = best_value
    # print(v_t)
    return v_t

v_t_star = value_iteration()
v_star = evaluate_values().T
print(v_t_star)
print(v_star)
print("State value function in part a and b are {}".format("the same" if np.all((abs(
v_t_star - v_star)<1e-2)) else "different"))
```

```
[ 0. 0. 0. 0. 53.50369907
 0. 0. 0. 0. 0.
 65.76590103 55.87677976 54.91691772 54.13958365 52.49819804
 43.76763184 47.94774456 0. 0. 67.12923121
-99.99015743 50.47110082 0. 43.93106347 -99.99015743
 48.76348165 0. 0. 77.83779421 70.30077743
 59.66008697 -99.99015743 51.08596413 0. 58.14122806
 0. 0. 79.8361461 81.3359714 0.
-99.99015743 61.00082068 0. 59.38385357 0.
 0. 72.46759065 83.03997576 80.94995886 61.77343094
 71.77908872 70.34422529 60.16270949 0. 0.
-99.99015743 84.87195914 0. -99.99015743 73.93914581
 0. -99.99015743 0. 0. 0.
 96.86264728 97.03507372 88.21156951 85.17609259 -99.99015743
 0. 0. 0. 99.99015743 98.70902034
 98.71747665 99.99015743 97.56294932 88.3971411 99.99015743
 0. ]

[[ 0. 0. 0. 0. 53.50968756
 0. 0. 0. 0. 0.
 65.77308407 55.88294346 54.92298013 54.14557214 52.50402036
 43.77254574 47.95296148 0. 0. 67.13647421
-100. 50.47656297 0. 43.9359876 -100.
 48.76871928 0. 0. 77.84605 70.30818136
 59.66641187 -100. 51.09137525 0. 58.14735126
 0. 0. 79.84451583 81.34440225 0.
-100. 61.00715483 0. 59.39003194 0.
 0. 72.47511769 83.04847989 80.95826449 61.77980767
 71.78642614 70.35142939 60.1688947 0. 0.
-100. 84.88054612 0. -100. 73.94661407
 0. -100. 0. 0. 0.
 96.87232244 97.04482865 88.22035599 85.18458536 -100.
 0. 0. 0. 100. 98.71875987
 98.72729893 100. 97.57257319 88.40593622 100.
 0. ]]
```

State value function in part a and b are the same